

Graphs - Traveling Salesman Problem | Nearest Neighbor Heuristic


Traveling Salesman Problem: Let us consider a traveling salesman that wants to visit each of N given cities $c = 0, \dots, N-1$ exactly once, return back to the city he started from and perform this tour with a minimum cost. For each pair of distinct cities c and c' there is an associated traversal cost $g(c, c')$ which in turn corresponds to the weight of the edge of the respective graph representation. Furthermore, it is assumed that the graph is complete - that is that we can go directly from every city to every other city. The TSP problem is to find a tour, a visit order, that goes through every city only once and the sum of the associated edges cost is minimum.

Nearest Neighbor Heuristic: There are many solutions to the TSP problem and in fact - as discussed - this is a known hard problem of mathematics. An imperfect, approximate, but fairly simple solution is the so-called "Nearest Neighbor" Heuristic (NNH).

NNH starts from a partial tour (i.e., an ordered collection of distinct cities) and proceeds to construct a sequence of partial tours, adding to each partial tour a new city that a) does not close a cycle in the graph, and b) minimizes the cost of the new extended tour. More specifically, given a sequence $\{c_0, c_1, \dots, c_k\}$ of distinct cities, NNH adds a city c_{k+1} that minimizes the associated cost $g(c_k, c_{k+1})$ over all cities other than those already visited. This in turn forms the sequence $\{c_0, c_1, \dots, c_k, c_{k+1}\}$. Performing this process incrementally, NNH eventually identifies a sequence of N cities $\{c_0, c_1, \dots, c_{N-1}\}$ thus resulting in complete tour with total cost:

$$g(c_0, c_1) + \dots + g(c_{N-2}, c_{N-1}) + g(c_{N-1}, c_0)$$

The associated cost is not the optimal minimum but the other constraints of the TSP problem (e.g., no cycle) are satisfied. It is thus a valid solution. An associated visualization of the incremental methodology that NNH is approximating a solution of TSP is shown below.



Task: You are asked to implement a C++ implementation of NNH which considers an undirected and fully connected graph (e.g., with 4 vertices and edge weights you will specify) and returns the cost of the Nearest Neighbor Heuristic solution.

References:

- Bertsekas, D.P., 2019. Reinforcement learning and optimal control. Athena Scientific.
- Flood, M.M., 1956. The traveling-salesman problem. Operations research, 4(1), pp.61-75.
- <https://www.geeksforgeeks.org/traveling-salesman-problem-tsp-implementation/>
(previously discussed approach using evaluation of permutations)
- <http://www.martinbroadhurst.com/nearest-neighbour-algorithm-for-tsp-in-c.html>

Indicative Solution

```
#include <stdlib.h>

typedef struct {
    unsigned int first;
    unsigned int second;
    unsigned int weight;
} weighted_edge;

/* Check if the tour already contains an edge */
static unsigned int tour_contains(const weighted_edge *tour, unsigned int
t,
    const weighted_edge *edge)
{
    unsigned int contains = 0;
    unsigned int i;
    for (i = 0; i < t && !contains; i++) {
        contains = tour[i].first == edge->first
            && tour[i].second == edge->second;
    }
    return contains;
}

/* Find the edge to v's nearest neighbour not in the tour already */
static unsigned int nearest_neighbour_edge(const weighted_edge *edges,
unsigned int size,
    const weighted_edge *tour, unsigned int t, unsigned int v)
{
    unsigned int min_distance = 0;
    unsigned int nearest_neighbour;
    unsigned int i;
    for (i = 0; i < size; i++) {
        if ((edges[i].first == v || edges[i].second == v)
            && (min_distance == 0 || edges[i].weight < min_distance)
            && !tour_contains(tour, t, &edges[i]))
        {
            min_distance = edges[i].weight;
            nearest_neighbour = i;
        }
    }
}
```

```

    }
    return nearest_neighbour;
}

weighted_edge *nearest_neighbour_tsp(const weighted_edge *edges, unsigned
int size,
    unsigned int order)
{
    unsigned int t, v = 0;
    weighted_edge *tour = malloc(order * sizeof(weighted_edge));
    if (tour == NULL) {
        return NULL;
    }
    for (t = 0; t < order; t++) {
        unsigned int e = nearest_neighbour_edge(edges, size, tour, t, v);
        tour[t] = edges[e];
        v = edges[e].first == v ? edges[e].second : edges[e].first;
    }
    return tour;
}

#include <stdio.h>
#include <stdlib.h>
/* Connect two edges */
void weighted_edge_connect(weighted_edge *edges, unsigned int first,
unsigned int second,
    unsigned int weight, unsigned int *pos)
{
    edges[*pos].first = first;
    edges[*pos].second = second;
    edges[*pos].weight = weight;
    (*pos)++;
}

void print_edges(const weighted_edge *edges, unsigned int n)
{
    unsigned int e;

```

```

    for (e = 0; e < n; e++) {
        printf("(%u, %u, %u) ", edges[e].first, edges[e].second,
edges[e].weight);
    }
    putchar('\n');
}
int main(void)
{
    unsigned int i = 0;
    const unsigned int size = 15; /* Edges */
    const unsigned int order = 6; /* Vertices */
    weighted_edge *edges = malloc(size * sizeof(weighted_edge));
    weighted_edge *tour;
    weighted_edge_connect(edges, 0, 1, 25, &i);
    weighted_edge_connect(edges, 0, 2, 19, &i);
    weighted_edge_connect(edges, 0, 3, 19, &i);
    weighted_edge_connect(edges, 0, 4, 16, &i);
    weighted_edge_connect(edges, 0, 5, 28, &i);
    weighted_edge_connect(edges, 1, 2, 24, &i);
    weighted_edge_connect(edges, 1, 3, 30, &i);
    weighted_edge_connect(edges, 1, 4, 27, &i);
    weighted_edge_connect(edges, 1, 5, 17, &i);
    weighted_edge_connect(edges, 2, 3, 18, &i);
    weighted_edge_connect(edges, 2, 4, 20, &i);
    weighted_edge_connect(edges, 2, 5, 23, &i);
    weighted_edge_connect(edges, 3, 4, 19, &i);
    weighted_edge_connect(edges, 3, 5, 32, &i);
    weighted_edge_connect(edges, 4, 5, 41, &i);
    tour = nearest_neighbour_tsp(edges, size, order);
    print_edges(tour, order);

    free(tour);
    free(edges);
    return 0;
}

```

